

SQL: Data Manipulation Language

Part 1

Warning!

Do not confuse

the *relational algebra* **SELECT** operator
with the *SQL* **SELECT** operator...

they are not the same!

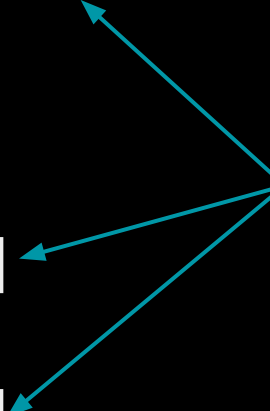
SELECT Syntax

```
SELECT [DISTINCT] column_list1  
      [INTO table2]  
FROM table_list  
[WHERE condition1]  
[GROUP BY column_list2]  
[HAVING condition2]  
[ORDER BY column_list3]
```

SELECT Syntax

```
SELECT [DISTINCT] column_list1
      [INTO table2]
      FROM table_list
      [WHERE condition1]
      [GROUP BY column_list2]
      [HAVING condition2]
      [ORDER BY column_list3]
```

column_list is one or more columns separated by commas

The diagram consists of three blue arrows originating from a single point on the right side of the text box. One arrow points to the *column_list1* placeholder in the SELECT clause. The second arrow points to the *column_list2* placeholder in the GROUP BY clause. The third arrow points to the *column_list3* placeholder in the ORDER BY clause.

SELECT Syntax

```
SELECT [DISTINCT] column_list1  
      [INTO table2]  
FROM table_list ←  
[WHERE condition1]  
[GROUP BY column_list2]  
[HAVING condition2]  
[ORDER BY column_list3]
```

table_list is one or more tables
separated by commas

Clauses

- **SELECT** displays columns
- **DISTINCT** compresses out duplicate rows
- **FROM** specifies which tables to use
- **WHERE** chooses rows to include
- **GROUP BY** reorganizes data into bands
- **HAVING** chooses groups to include
- **ORDER BY** sorts the output

SELECT Syntax (more formally)

Subselect General Form:

```
SELECT [all|distinct] expression {, expression}  
    FROM tablename [corr_name] {, tablename [corr_name]}  
    [WHERE search_condition]  
    [GROUP BY column {, column}]  
    [HAVING search_condition]
```

Full Select General Form:

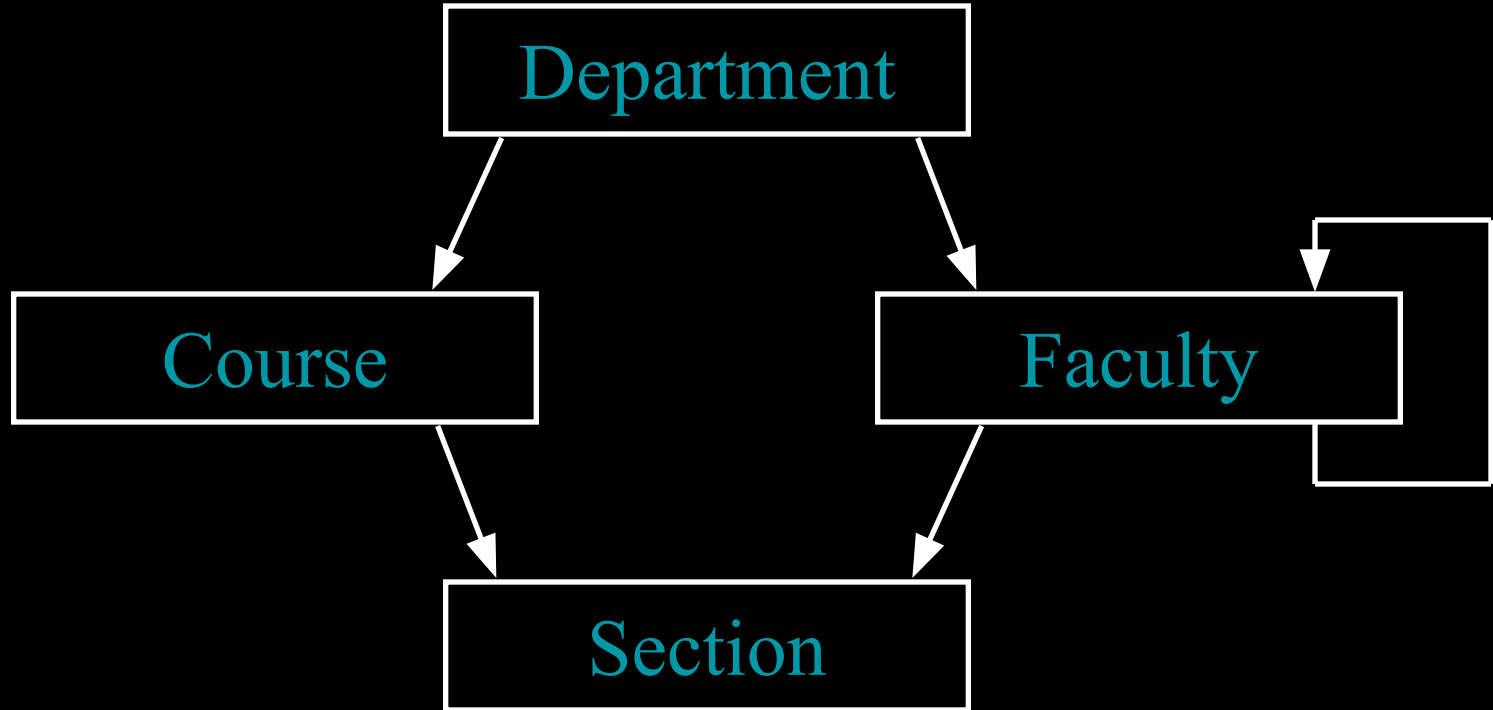
Subselect

```
{UNION [ALL] Subselect}  
[ORDER BY result_column [ ASC | DESC ] {, result_column [ ASC | DESC ]}]
```

Conceptual Order of Evaluation of a Select Statement

1. First the product of all tables in the **from** clause is formed.
2. The **where** clause is then evaluated to eliminate rows that do not satisfy the *search_condition*.
3. Next, the rows are grouped using the columns in the **group by** clause.
4. Then, Groups that do not satisfy the *search_condition* in the **having clause** are eliminated.
5. Next, the expressions in the **select** clause target list are evaluated.
6. If the **distinct** keyword is present in the select clause, duplicate rows are now eliminated.
7. The **union** is taken after each sub-select is evaluated.
8. Finally, the resulting rows are sorted according to the columns specified in the **order by** clause.

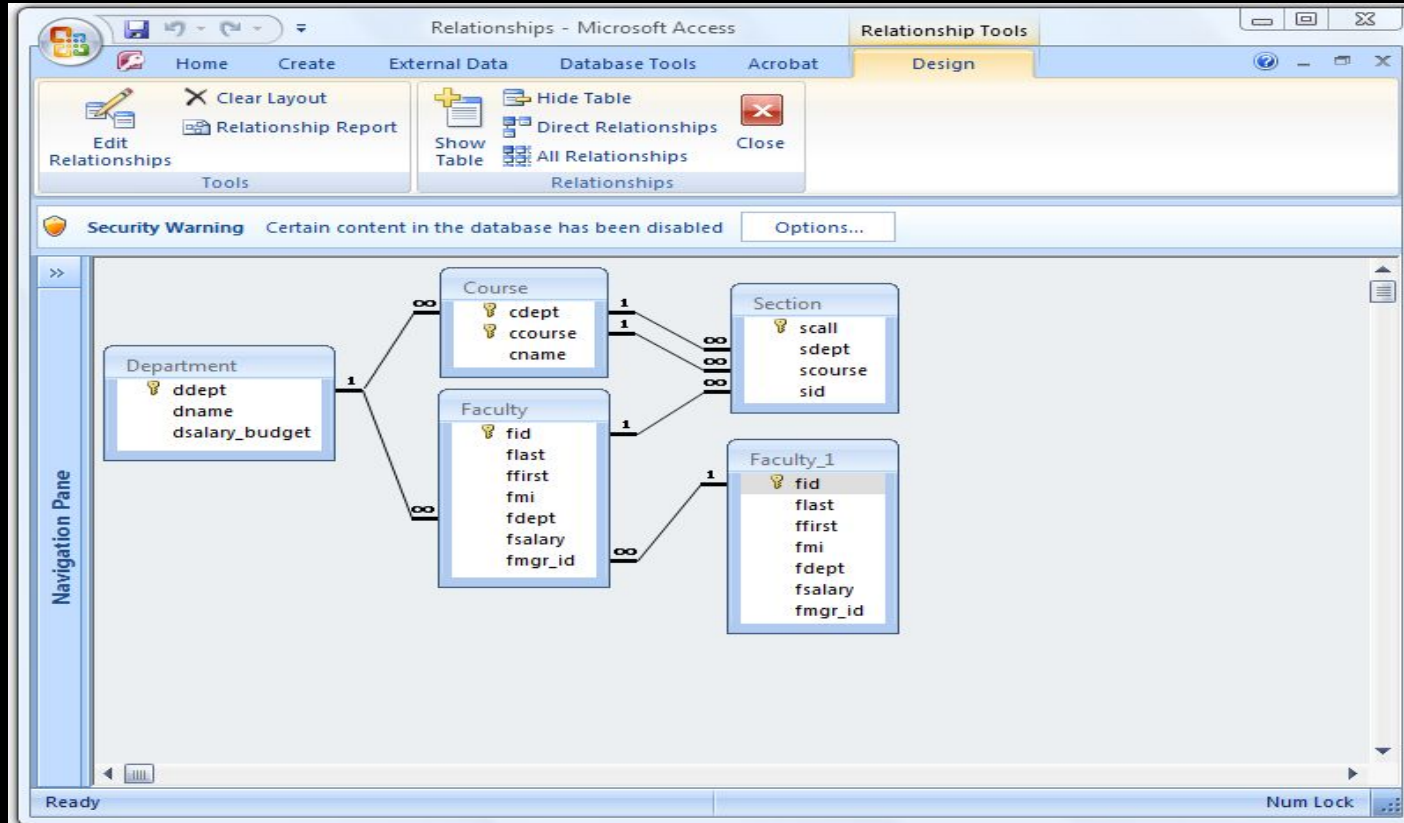
Sample Database



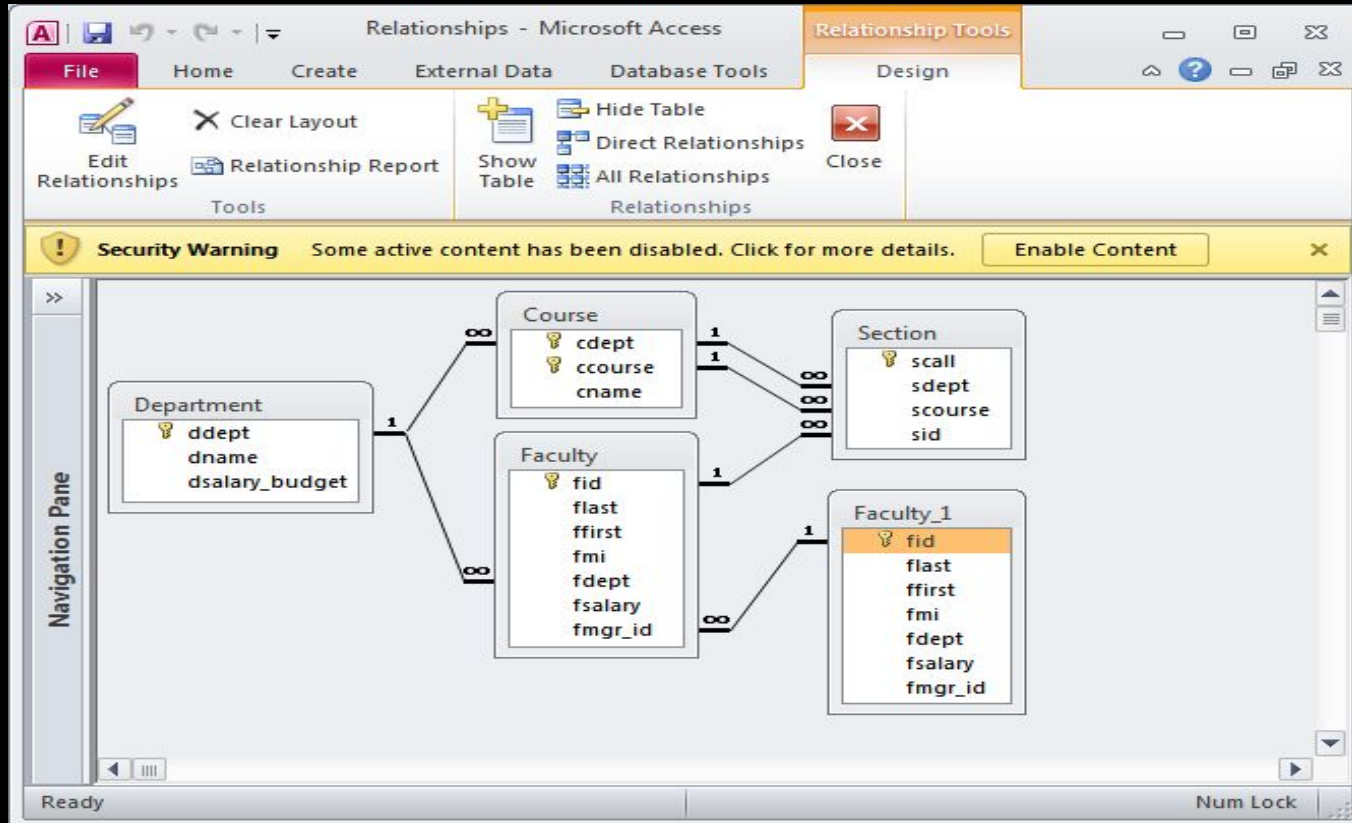
Relationships

Table	Primary Key	Foreign Key
Department	ddept	
Course	cdept & ccourse	cdept
Faculty	fid	fdept fmgr_id
Section	scall	sid sdept & scourse

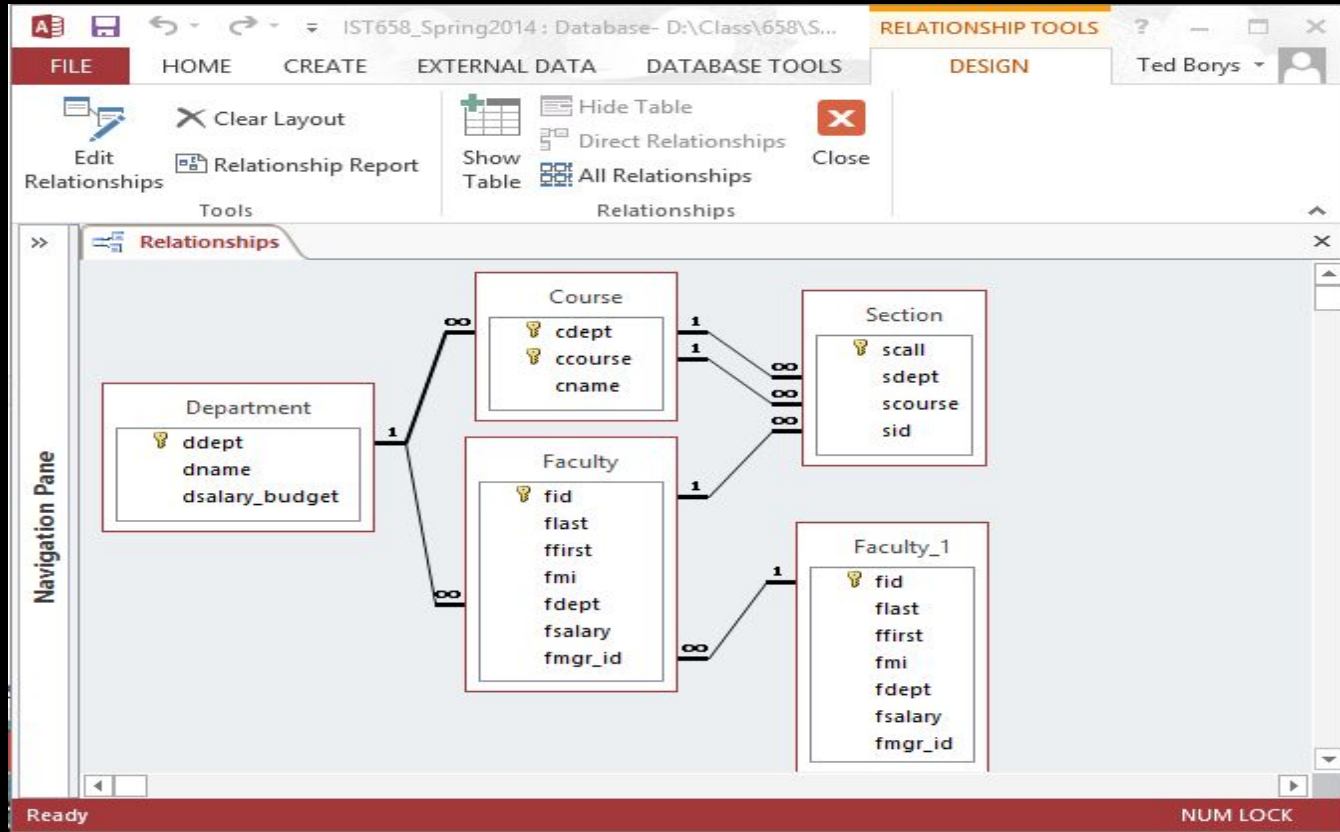
Access 2007



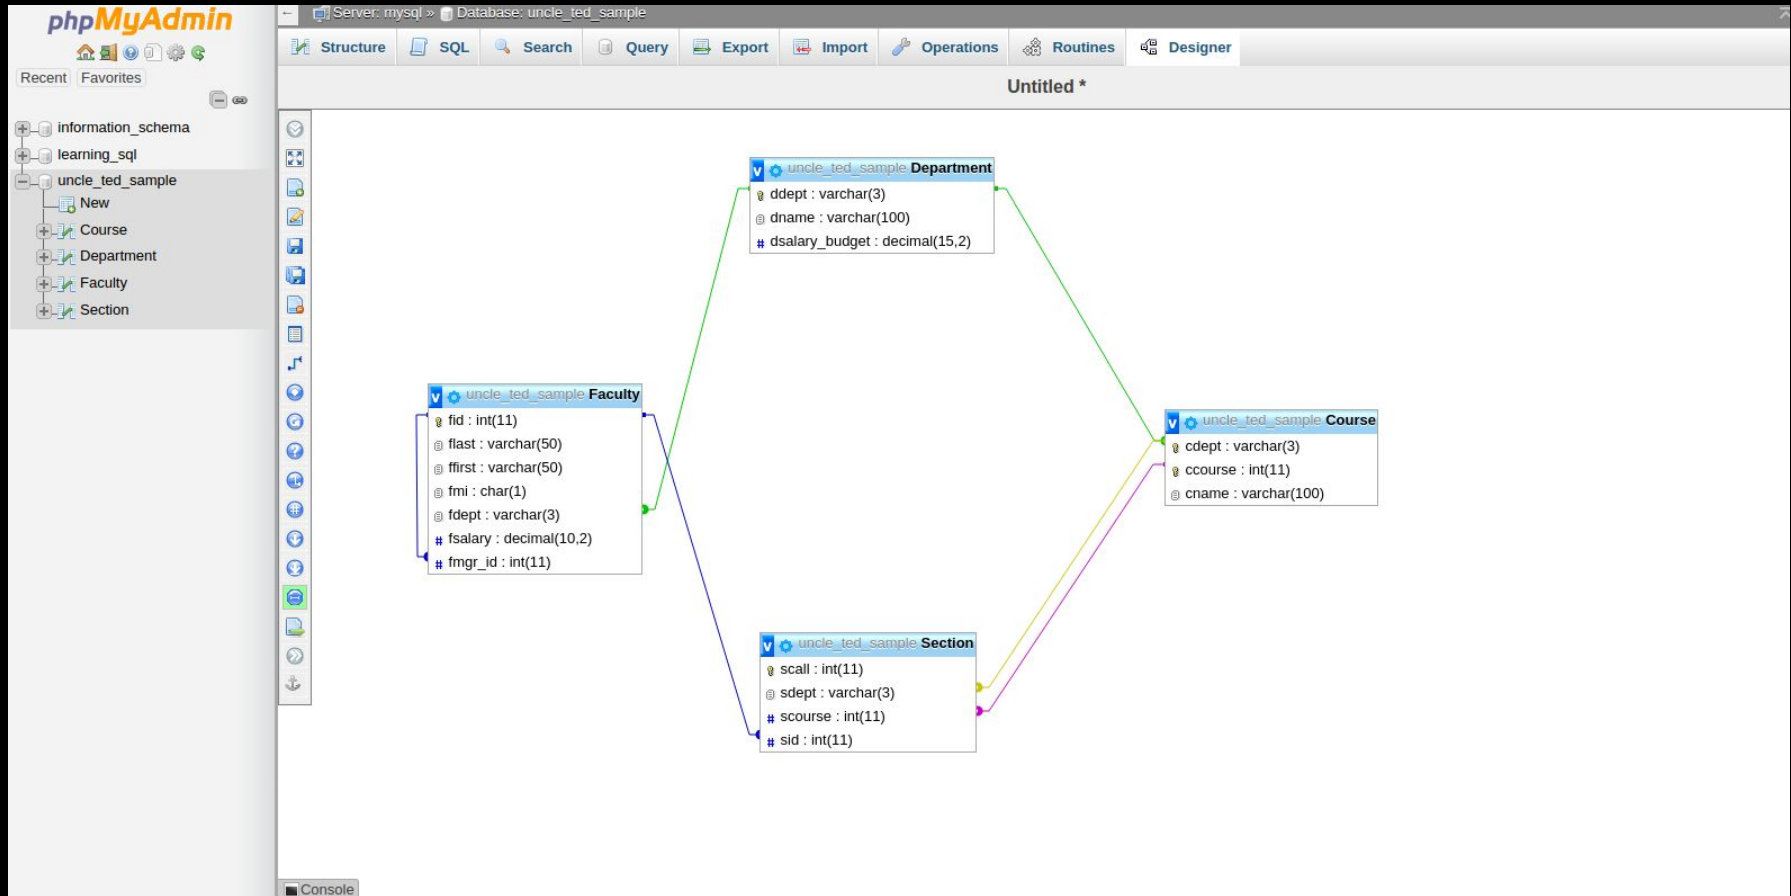
Access 2010



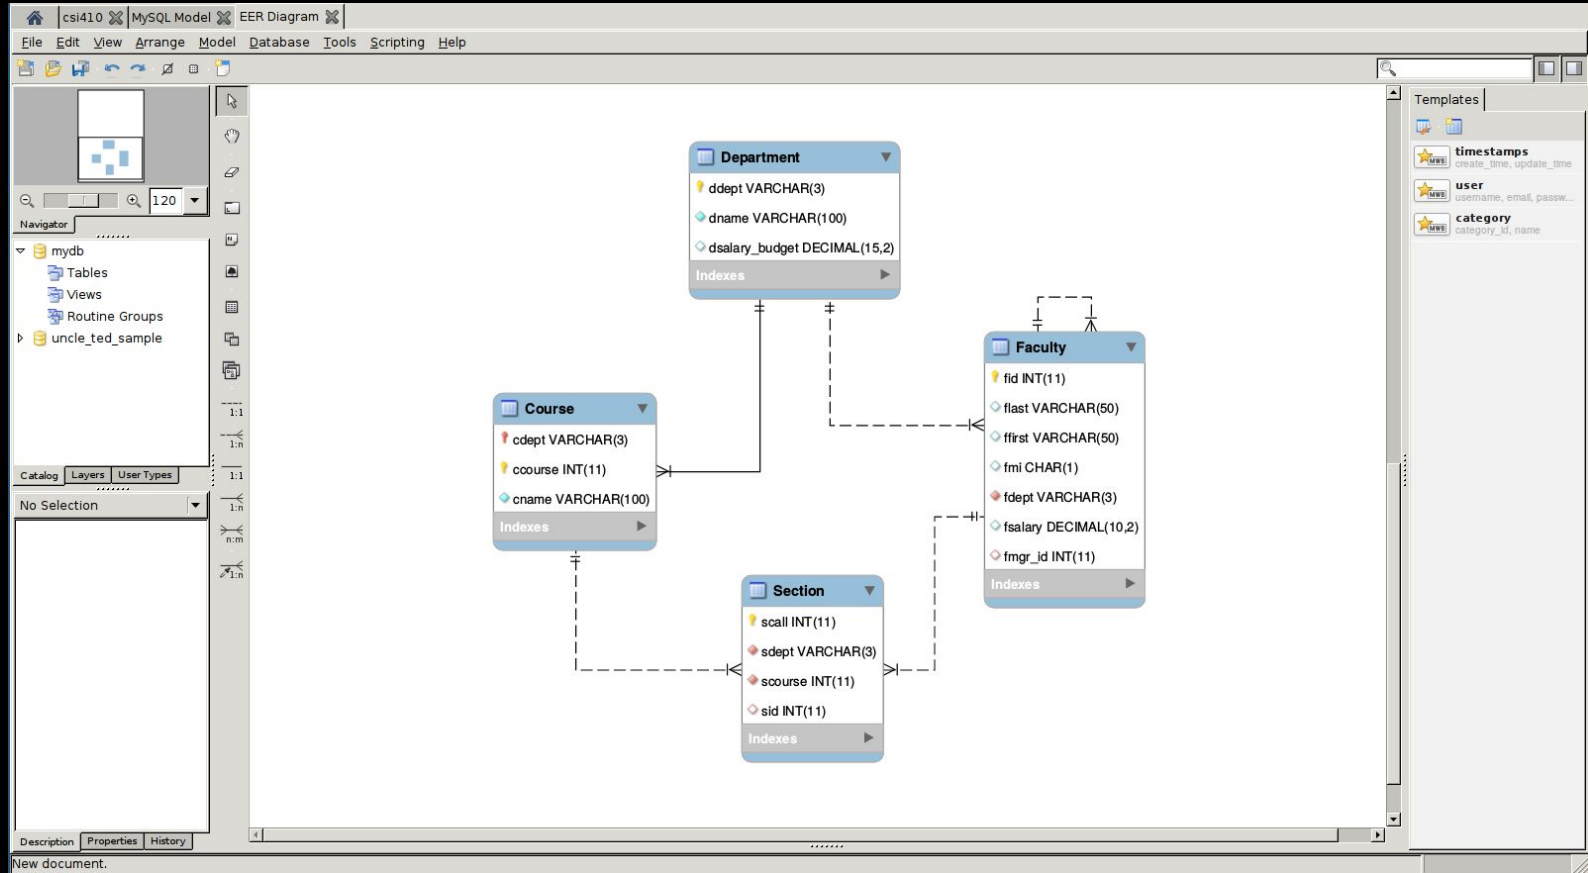
Access 2013



phpMySQL



MySQL Workbench



Department

<u>ddept</u>	dname	dsalary_budget
ADM	Administration	200000
ATM	Atmospheric Science	90000
BIO	Biology	40000
CSI	Computer Science	180000
ENG	English	80000
SPN	Spanish	70000

Faculty

<u>fid</u>	flast	ffirst	fmi	fdept	fsalary	fmgr_id
12058	Borys	Ted	J	CSI	48000	22321
12206	Ryan	Alfred	C	ENG	48000	52110
21004	Perry	Bill	S	BIO	21800	31890
22321	Brady	Kathy	M	CSI	63400	52110
31890	Coulsen	Mary		BIO	21400	52110
32000	delBene	Bill	S	CSI	63500	22321
47862	Anders	John	P	ENG	33700	12206
52110	Smith	Alice		ADM	82000	<i>null</i>

Course

<u>cdept</u>	<u>ccourse</u>	cname
ATM	408	Hydrometeorology
ATM	410	Dynamic Meteorology 1
BIO	205	Human Genetics
BIO	410	Human Physiology
CSI	205	C Language Programming
CSI	409	Automata & Formal Languages
CSI	410	Database Management Systems
ENG	427	The Victorian Period

Section

<u>sclass</u>	sdept	scourse	sid
102	ATM	408	22321
273	BIO	205	21004
285	BIO	410	22321
312	CSI	205	22321
313	CSI	205	47862
324	CSI	410	12058

Query 1

Display all the data in the Department table.

Query 1

Display all the data in the Department table.

```
SELECT ddept, dname, dsalary_budget  
FROM   Department;
```

Query 1

Display all the data in the Department table

```
SELECT ddept, dname, dsalary_budget  
FROM   Department;
```

or

```
SELECT *  
FROM   Department;
```

Query 1

Display all the data in the Department table

```
SELECT ddept, dname, dsalary_budget  
FROM   Department;
```

or

```
SELECT *  
FROM Department;
```

or

```
SELECT Department.*  
FROM Department;
```

Query 2

Display the flast, ffirst, fdept, and fsalary columns from the Faculty table.

Query 2

Display the flast, ffirst, fdept, and fsalary columns from the Faculty table.

```
SELECT flast, ffirst, fdept, fsalary  
FROM Faculty;
```

Query 3

Display all the first names in the Faculty table and remove duplicate rows.

Query 3

Display all the first names in the Faculty table and remove duplicate rows.

```
SELECT DISTINCT ffirst  
FROM Faculty;
```

SQL Seduction 1

DISTINCT compresses out fully duplicated rows, i.e., it is applied to the combination of all the values of all the columns being displayed, not just the first column.

```
SELECT DISTINCT ffirst, flast  
FROM Faculty;
```

Query 4

Display all the data in the Course table in *descending cname order*.

Query 4

Display all the data in the Course table in *descending cname order*.

```
SELECT *  
FROM Course  
ORDER BY cname DESC;
```

Query 5

Display the id, full name, and salary columns from the Faculty table and sort the output in telephone book order.

Query 5

Display the id, full name, and salary columns from the Faculty table and sort the output in telephone book order.

```
SELECT fid, ffirst, fmi, flast, fsalary  
FROM Faculty  
ORDER BY flast, ffirst, fmi;
```


How to Join Tables

- Follow the relationships
 - Equate owner's primary key to member's foreign key
 - Equate the primary key of one table to the primary key of another table (1:1 relationship)
- Intelligently joining on foreign key pairs is not likely
- Fight temptation to join on any pair of columns that appear similar
- Of course, there are exceptions ...

Query 6

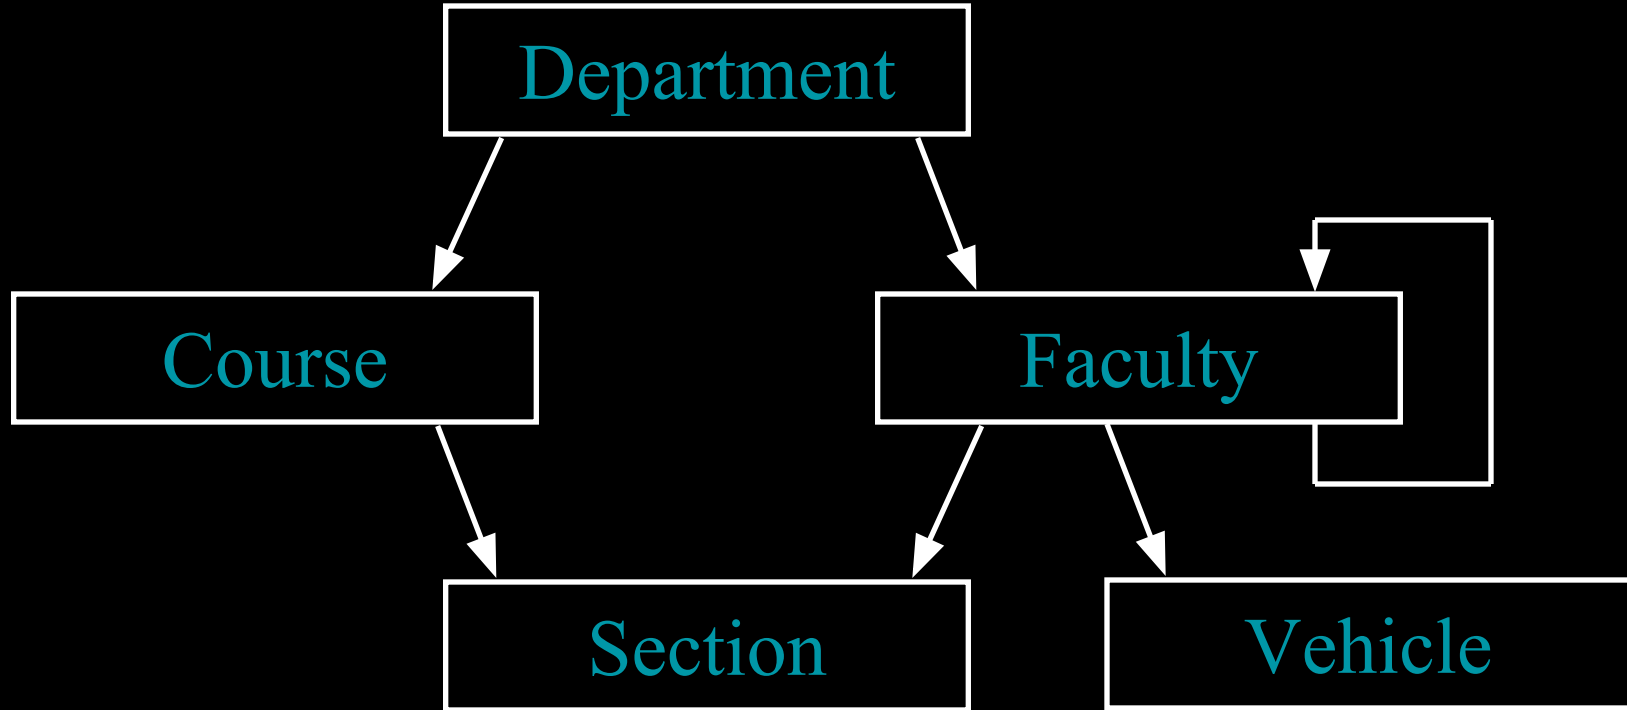
Join the Department and Faculty tables and display all the columns.

Query 6

Join the Department and Faculty tables and display all the columns.

```
SELECT *  
FROM Department, Faculty  
WHERE ddept = fdept;
```

Modified Sample Database



Continuing The Set Up...

- Eye color column added to Faculty: feye
- Vehicle table added
 - Primary key is vpermit#
 - Foreign key is vid
 - Vehicle has a color column: vcolor

SQL Seduction 2

Display all information from the Faculty and Vehicle tables for faculty whose eye color matches their vehicle's color.

SQL Seduction 2

Display all information from the Faculty and Vehicle tables for faculty whose eye color matches their vehicle's color.

```
SELECT *  
FROM Faculty, Vehicle  
WHERE feye = vcolor;
```

Gotcha!

SQL Seduction 2

Display all information from the Faculty and Vehicle tables for faculty whose eye color matches their vehicle's color.

```
SELECT *  
FROM Faculty, Vehicle  
WHERE (feye = vcolor) AND (fid = vid);
```

We need to use the foreign key!!!



Query 7

Join the Faculty and Course tables.

Query 7

Join the Faculty and Course tables.

```
SELECT *  
FROM Faculty, Course  
WHERE fdept = cdept;
```

**Gotcha
again!**

Query 7

Join the Faculty and Course tables.

```
SELECT Faculty.*, Course.*  
FROM Faculty, Course, Section  
WHERE (fid = sid)  
      AND (cdept = fdept)  
      AND (ccourse = scourse);
```

ANOTHER
GOTCHA!!!

What about instructors
teaching outside of their
departments?

Query 7

Join the Faculty and Course tables.

```
SELECT Faculty.*, Course.*  
FROM Faculty, Course, Section  
WHERE (fid = sid)  
      AND (cdept = sdept)  
      AND (ccourse = scourse);
```

Query 7

This query shows us the rows missed when we assume instructors only teach within their dept.

```
SELECT Faculty.*, Course.*  
FROM Faculty, Course, Section  
WHERE (fid = sid)  
      AND (ccourse = scourse)  
      AND (cdept = sdept)  
      AND (fdept <> cdept);
```

fid	flast	ffirst	fmi	fdept	fsalary	fmgr_id	cdept	ccourse	cname
22321	Brady	Kathy	M	CSI	63400	52110	ATM	408	Hydrometeorology
22321	Brady	Kathy	M	CSI	63400	52110	BIO	410	Human Physiology
47862	Anders	John	P	ENG	33700	12206	CSI	205	C Language Programming

NOTE: The table name prefixes for the column names were removed.

Query 8

Join the Department and Course tables.

Query 8

Join the Department and Course tables.

```
SELECT *  
FROM Department, Course  
WHERE ddept = cdept;
```

Query 9

Join the Course and Section tables.

Query 9

Join the Course and Section tables.

```
SELECT *  
FROM Course, Section  
WHERE (cdept=sdept)  
      AND (ccourse=scourse);
```

Query 10

Do a 3-way join of the Department, Course, and Section tables.

Query 10

Do a 3-way join of the Department, Course, and Section tables.

```
SELECT *  
FROM Department, Course, Section  
WHERE (ddept = cdept)  
      AND (cdept = sdept)  
      AND (ccourse = scourse);
```

Query 11

Do a 3-way join of the Department, Faculty, and Section tables.

Query 11

Do a 3-way join of the Department, Faculty, and Section tables.

```
SELECT *  
FROM Department, Faculty, Section  
WHERE (ddept=fdept) AND (fid = sid);
```

Query 12

Incorrectly join the Department, Faculty, and Section tables (another SQL Seduction 2)

```
SELECT *  
FROM Department, Faculty, Section  
WHERE ddept=fdept AND fdept=sdept;
```

Query 12

Incorrectly join the Department, Faculty, and Section tables (another SQL Seduction 2)

```
SELECT *  
FROM Department, Faculty, Section  
WHERE (ddept = fdept)  
      AND (fid = sid)  
      AND (fdept = sdept);
```

Typical Subqueries

- **WHERE** *column_name* **IN** (*subquery*)
- **WHERE EXISTS** (*subquery*)
- **Standalone**
 - Can be executed by itself
 - Typically used with **IN**
- **Correlated**
 - Subquery uses column(s) from outer query's table
 - Can't be executed by itself
 - Typically used with **EXISTS**

Query 13

Find departments that offer courses.

Query 13

Find departments that offer courses.

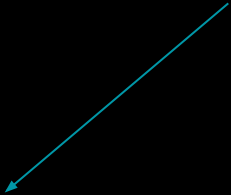
```
SELECT *  
FROM Department  
WHERE ddept IN  
    (SELECT cdept FROM Course);
```

How Standalone Subquery Works

- Subquery, inside parentheses, executes
 - Intermediate result generated once
 - Original statement reduces to:

```
SELECT *  
FROM Department  
WHERE ddept IN  
    ("ATM","ATM","BIO","BIO","CSI","CSI","CSI","ENG");
```

(SELECT cdept FROM Course)



How Standalone Subquery Works

- May want to use **DISTINCT** in subquery
 - This is only a potential performance issue
 - Correct result generated either way

Query 14

Another way to find departments that offer courses.

Query 14

Another way to find departments that offer courses.

```
SELECT *  
FROM Department  
WHERE EXISTS (  
    SELECT cdept  
    FROM Course  
    WHERE cdept = ddept);
```

How Correlated Subquery Works

- **EXISTS** checks to see if rows are generated in subquery for each row examined in the main query.
- Subquery, inside parentheses, executes
 - Intermediate result generated for each row in Department.

Examining 1st Department Row

... the original statement internally reduces to:

```
SELECT *  
FROM Department  
WHERE EXISTS  
  (SELECT cdept  
   FROM Course  
   WHERE cdept = "ADM" );
```



```
SELECT *  
FROM Department  
WHERE EXISTS  
  ();
```


Examining 1st Department Row

... the original statement internally reduces to:

```
SELECT *  
FROM Department  
WHERE EXISTS  
  (SELECT cdept  
   FROM Course  
   WHERE cdept = "ADM" );
```



```
SELECT *  
FROM Department  
WHERE EXISTS  
  ("ATM", "ATM");
```

IN versus EXISTS

- All queries that use **IN** can be re-written with **EXISTS**.
- Some queries that use **EXISTS** can be re-written with **IN**, **but not all can.**
- The theoretical reason why is beyond the scope of this class.

Query 15

And one more way to find departments that offer courses.

```
SELECT DISTINCT Department.*  
FROM Department, Course  
WHERE ddept = cdept;
```

Query 16

You can correlate this subquery. Does no logical harm, but adds *syntax clutter*.

```
SELECT *  
FROM Department  
WHERE ddept IN  
  (SELECT cdept  
   FROM Course  
   WHERE cdept = ddept);
```

Query 17

Even this variation works correctly, despite the added *syntax clutter*.

```
SELECT *  
FROM Department  
WHERE ddept IN  
    (SELECT cdept  
     FROM Course, Department  
     WHERE cdept = ddept);
```

SQL Seduction 3

But, this variation does not work. As long as there is at least one row in Course, every Department row will be displayed!

```
SELECT *  
FROM Department  
WHERE EXISTS  
(SELECT cdept  
FROM Course, Department  
WHERE cdept = ddept);
```

Query 18

Find departments that don't offer courses.

Query 18

Find departments that don't offer courses.

```
SELECT *  
FROM Department  
WHERE ddept NOT IN  
      (SELECT cdept FROM Course);
```


Query 19

Another way to find departments that don't offer courses.


```
SELECT *  
FROM Department  
WHERE NOT EXISTS  
    (SELECT cdept  
     FROM Course  
     WHERE cdept = ddept);
```

SQL Seduction 4

Find departments that don't offer courses.

But this dog won't hunt.

You can't rewrite **NOT IN** as not equal.



```
SELECT DISTINCT Department.*  
FROM Department, Course  
WHERE ddept <> cdept;
```

Query 20

Find courses that have sections.

Query 20

Find courses that have sections.

```
SELECT DISTINCT Course.*  
FROM Course, Section  
WHERE (cdept = sdept)  
      AND (ccourse = scourse);
```

Query 21

Another way to find courses that have sections.

```
SELECT *  
FROM Course  
WHERE cdept&ccourse IN  
  (SELECT sdept&scourse FROM Section);
```

Use Access concatenation operator & to create single column.

Query 22

SQL99 standard lets you rewrite Query 21 like this, but Access 2007 disallows it.

```
SELECT *  
FROM Course  
WHERE (cdept, ccourse) IN  
      (SELECT sdept, scourse FROM Section);
```

SQL Seduction 5

You can't replace the single, concatenated IN with two, singular INs (Hint: ATM 410).

```
SELECT *  
FROM Course  
WHERE cdept IN (SELECT sdept FROM Section)  
AND ccourse IN(SELECT scourse FROM Section);
```

Query 23

And one more way to find courses that have sections

```
SELECT *  
FROM Course  
WHERE EXISTS  
  (SELECT * FROM Section  
    WHERE (cdept = sdept)  
          AND (ccourse = scourse));
```


Query 24

This is another incarnation of SQL Seduction 5

```
SELECT * FROM Course  
WHERE EXISTS  
  (SELECT * FROM Section WHERE cdept = sdept)  
AND EXISTS  
  (SELECT * FROM Section WHERE ccourse = scourse);
```

Table Name Prefix

- Must be used when column names are not unique based on tables named in **FROM** clause
- Otherwise, can be used anytime
 - IMHO, adds syntax clutter

```
SELECT Course.cdept, Course.ccourse, Course.cname  
FROM Course  
ORDER BY Course.cname;
```

Table Name Alias

- Must be used when joining a table back on itself (also known as a *self-join*)
- Otherwise, can be used anytime
 - Useful when table name prefixing required, and table names are long
 - IMHO, when not needed, adds syntax clutter

```
SELECT C.cdept, C.ccourse, C.cname  
FROM Course C  
ORDER BY C.cname;
```

Query 25

Join subordinates and managers... they both just happen to be in the same table: Faculty.

Query 25

Join subordinates and managers... they both just happen to be in the same table: Faculty.

```
SELECT *  
FROM Faculty Sub, Faculty Mgr  
WHERE Sub.fmgr_id = Mgr.fid;
```

Query 26

Find all the subordinates who make more than their managers.

Query 26

Find all the subordinates who make more than their managers.

```
SELECT *  
FROM Faculty Sub, Faculty Mgr  
WHERE (Sub.fmgr_id = Mgr.fid)  
      AND (Sub.fsalary > Mgr.fsalary);
```

Query 27

List all the departments in Faculty and Course.

Query 27

List all the departments in Faculty and Course.

```
SELECT fdept FROM Faculty  
UNION  
SELECT cdept FROM Course;
```

Query 27

List all the departments in Faculty and Course.

```
SELECT fdept FROM Faculty  
UNION  
SELECT cdept FROM Course;
```

Note: duplicate rows eliminated

UNION ALL preserves duplicates

Query 28

Count the number of rows in Faculty.

Query 28

Count the number of rows in Faculty.

```
SELECT COUNT(*)  
FROM Faculty;
```

Query 29

Count the number of *non-null* fid values in Faculty.

Query 29

Count the number of *non-null* fid values in Faculty.

```
SELECT COUNT(fid)
FROM Faculty;
```

Since fid is
primary key,
never null.

SQL Seduction 6

Counts non-null middle initial values; not necessarily the same as the number of rows.

```
SELECT COUNT(fmi)  
FROM Faculty;
```

Query 30

Count the number of departments that offer courses.

```
SELECT COUNT(DISTINCT cdept)  
FROM Course;
```

This syntax disallowed by Access.