

# SQL: Data Manipulation Language

Part 2

SQL Arithmetic Operators	
Operator	Description
+	Addition - Adds values on either side of the operator
-	Subtraction - Subtracts right hand operand from left hand operand
*	Multiplication - Multiplies values on either side of the operator
/	Division - Divides left hand operand by right hand operand
%	Modulus - Divides left hand operand by right hand operand and returns remainder

## SQL Comparison Operators

Operator	Description
=	Checks if the values of two operands are equal or not, if yes then condition becomes true.
!=	Checks if the values of two operands are equal or not, if values are not equal then condition becomes true.
<>	Checks if the values of two operands are equal or not, if values are not equal then condition becomes true.
>	Checks if the value of left operand is greater than the value of right operand, if yes then condition becomes true.
<	Checks if the value of left operand is less than the value of right operand, if yes then condition becomes true.
>=	Checks if the value of left operand is greater than or equal to the value of right operand, if yes then condition becomes true.
<=	Checks if the value of left operand is less than or equal to the value of right operand, if yes then condition becomes true.

# SQL Logical Operators

Operator	Description
<b>ALL</b>	The ALL operator is used to compare a value to all values in another value set.
<b>AND</b>	The AND operator allows the existence of multiple conditions in an SQL statement's WHERE clause.
<b>ANY</b>	The ANY operator is used to compare a value to any applicable value in the list according to the condition.
<b>BETWEEN</b>	The BETWEEN operator is used to search for values that are within a set of values, given the minimum value and the maximum value.
<b>EXISTS</b>	The EXISTS operator is used to search for the presence of a row in a specified table that meets certain criteria.
<b>IN</b>	The IN operator is used to compare a value to a list of literal values that have been specified.
<b>LIKE</b>	The LIKE operator is used to compare a value to similar values using wildcard operators.
<b>NOT</b>	The NOT operator reverses the meaning of the logical operator with which it is used. Eg: NOT EXISTS, NOT BETWEEN, NOT IN, etc. <b>This is a negate operator.</b>
<b>OR</b>	The OR operator is used to combine multiple conditions in an SQL statement's WHERE clause.
<b>IS NULL</b>	The NULL operator is used to compare a value with a NULL value.
<b>UNIQUE*</b>	The UNIQUE operator searches every row of a specified table for uniqueness (no duplicates). *Not

# Logic

- Two-valued logic, aka 2VL (George Boole)
  - True
  - False
- SQL uses 3VL (Jan Lukasiewicz)
  - True
  - Unknown
  - False

# 2VL Truth Tables

AND	True	False
True	True	False
False	False	False

OR	True	False
True	True	True
False	True	False

NOT	True	False
	False	True

# 3VL Truth Tables

AND	True	Unk	False
True	True	Unk	False
Unk	Unk	Unk	False
False	False	False	False

OR	True	Unk	False
True	True	True	True
Unk	True	Unk	Unk
False	True	Unk	False

NOT	True	Unk	False
	False	Unk	True

# NULL

## NOTE:

The following queries were run in PostgreSQL 9.5.

*This syntax will not work in MySQL.*

Additionally, the following psql shell command was used to change how NULLs are displayed: `\pset null '[NULL]'`



# NULL

- `SELECT true AND NULL;`
- `SELECT true OR NULL;`
- `SELECT false AND NULL;`
- `SELECT false OR NULL;`

# NULL

- `SELECT true AND NULL;`

?column?
[NULL]

- `SELECT false AND NULL;`

- `SELECT true OR NULL;`

- `SELECT false OR NULL;`

# NULL

- `SELECT true AND NULL;`

?column?
[NULL]

- `SELECT true OR NULL;`

- `SELECT false AND NULL;`

?column?
f

- `SELECT false OR NULL;`

# NULL

- `SELECT true AND NULL;`

?column?
[NULL]

- `SELECT false AND NULL;`

?column?
f

- `SELECT true OR NULL;`

?column?
t

- `SELECT false OR NULL;`

# NULL

- **SELECT** true AND NULL;

?column?
[NULL]

- **SELECT** true OR NULL;

?column?
t

- **SELECT** false AND NULL;

?column?
f

- **SELECT** false OR NULL;

?column?
[NULL]

# NULL

- `SELECT NOT NULL;`

# NULL

- `SELECT NOT NULL;`

?column?
[NULL]

# NULL

- `SELECT 1 WHERE 1 = 1;`
- `SELECT 1 WHERE 1 = NULL;`
- `SELECT 1 WHERE 1 = 0;`
- `SELECT 1 WHERE NULL = NULL;`

The WHERE will return a row only if the condition evaluates to true.  
UNKNOWN isn't true



# NULL

- `SELECT 1 WHERE 1 = 1;`

?column?
1

- `SELECT 1 WHERE 1 = NULL;`

- `SELECT 1 WHERE 1 = 0;`

- `SELECT 1 WHERE NULL = NULL;`

The WHERE will return a row only if the condition evaluates to true.  
UNKNOWN isn't true

# NULL

- `SELECT 1 WHERE 1 = 1;`

?column?
1

- `SELECT 1 WHERE 1 = NULL;`

- `SELECT 1 WHERE 1 = 0;`

?column?
----------

- `SELECT 1 WHERE NULL = NULL;`

The WHERE will return a row only if the condition evaluates to true.  
UNKNOWN isn't true

# NULL

- `SELECT 1 WHERE 1 = 1;`

?column?
1

- `SELECT 1 WHERE 1 = NULL;`

?column?
----------

- `SELECT 1 WHERE 1 = 0;`

?column?
----------

- `SELECT 1 WHERE NULL = NULL;`

The WHERE will return a row only if the condition evaluates to true.  
UNKNOWN isn't true

# NULL

- `SELECT 1 WHERE 1 = 1;`

?column?
1

- `SELECT 1 WHERE 1 = NULL;`

?column?
----------

- `SELECT 1 WHERE 1 = 0;`

?column?
----------

- `SELECT 1 WHERE NULL = NULL;`

?column?
----------

The WHERE will return a row only if the condition evaluates to true.  
UNKNOWN isn't true

# NULL

- `SELECT NULL = NULL;`
- `SELECT NULL != NULL;`

# NULL

- `SELECT NULL = NULL;`

?column?
[NULL]

- `SELECT NULL != NULL;`

# NULL

- `SELECT NULL = NULL;`

?column?
[NULL]

- `SELECT NULL != NULL;`

?column?
[NULL]

# NULL

How do we test for  
NULL?

- `SELECT NULL = NULL;`

?column?
[NULL]

- `SELECT NULL != NULL;`

?column?
[NULL]



# NULL

## The IS NULL operator .

- `SELECT 1 IS NULL;`

?column?
f

- `SELECT NULL IS NULL;`

?column?
t

- `SELECT 1 IS NOT NULL;`

?column?
t

- `SELECT NULL IS NOT NULL;`

?column?
f

# NULL

- `SELECT 1 + NULL;`

?column?
[NULL]

- `SELECT 1 < NULL;`

?column?
[NULL]

- `SELECT 1 > NULL;`

?column?
[NULL]

- `SELECT 1 <> NULL;`

?column?
[NULL]

# NULL

- `SELECT 0 BETWEEN 0 AND 2;`

?column?
t

- `SELECT 1 BETWEEN 0 AND 2;`

?column?
t

- `SELECT 0 BETWEEN 0 AND -1;`

?column?
f

# NULL

- `SELECT 0 BETWEEN 0 AND NULL;`

# NULL

- `SELECT 0 BETWEEN 0 AND NULL;`

?column?
[NULL]

# IN vs EXISTS vs JOIN

- **IN:**
  - Returns true if a specified value matches any value in a subquery or a list.
- **EXIST:**
  - Returns true if a subquery contains any rows.
- **JOIN:**
  - Joins 2 result sets on the joining

# NULL

- `SELECT EXISTS(SELECT NULL);`
- `SELECT NOT EXISTS(SELECT NULL);`

# NULL

- `SELECT EXISTS(SELECT NULL);`

?column?
t

- `SELECT NOT EXISTS(SELECT NULL);`



# NULL

- `SELECT EXISTS(SELECT NULL);`

?column?
t

- `SELECT NOT EXISTS(SELECT NULL);`

?column?
f

# NULL

- `SELECT 1 IN (1);`
- `SELECT 1 IN (NULL);`
- `SELECT 1 NOT IN (1);`
- `SELECT 1 NOT IN (NULL);`

# NULL

- `SELECT 1 IN (1);`

?column?
t

- `SELECT 1 IN (NULL);`

- `SELECT 1 NOT IN (1);`

- `SELECT 1 NOT IN (NULL);`

# NULL

- `SELECT 1 IN (1);`

?column?
t

- `SELECT 1 IN (NULL);`

- `SELECT 1 NOT IN (1);`

?column?
f

- `SELECT 1 NOT IN (NULL);`

# NULL

- `SELECT 1 IN (1);`

?column?
t

- `SELECT 1 IN (NULL);`

?column?
[NULL]

- `SELECT 1 NOT IN (1);`

?column?
f

- `SELECT 1 NOT IN (NULL);`

# NULL

- `SELECT 1 IN (1);`

?column?
t

- `SELECT 1 IN (NULL);`

?column?
[NULL]

- `SELECT 1 NOT IN (1);`

?column?
f

- `SELECT 1 NOT IN (NULL);`

?column?
[NULL]

# NULL

```
SELECT *  
FROM (  
    VALUES (NULL), (NULL)  
) AS T;
```

# NULL

```
SELECT *  
FROM (  
    VALUES (NULL), (NULL)  
) AS T;
```

column1
[NULL]
[NULL]



# NULL

```
SELECT DISTINCT *  
FROM (  
    VALUES (NULL), (NULL)  
) AS T;
```

# NULL

```
SELECT DISTINCT *  
FROM (  
    VALUES (NULL), (NULL)  
) AS T;
```

column1
[NULL]

# NULL

```
DROP TABLE IF EXISTS T1;  
DROP TABLE IF EXISTS T2;
```

```
CREATE TABLE t1 (  
  a INT  
);
```

```
INSERT INTO t1 (a)  
VALUES (1), (2),(null);
```

```
CREATE TABLE t2 AS  
SELECT * FROM t1;
```

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT *  
FROM t1, t2  
WHERE t1.a = t2.a;
```

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT *  
FROM t1, t2  
WHERE t1.a = t2.a;
```

a	a
1	1
2	2

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT *  
FROM t1 NATURAL JOIN t2;
```

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT *  
FROM t1 NATURAL JOIN t2;
```

t1	t2
a	a
1	1
2	2
null	null

a	a
1	1
2	2

# NULL

```
SELECT * FROM t1  
UNION ALL  
SELECT * FROM t2;
```

t1	t2
a	a
1	1
2	2
null	null



# NULL

```
SELECT * FROM t1  
UNION ALL  
SELECT * FROM t2;
```

a
1
2
null
1
2
null

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT * FROM t1  
UNION  
SELECT * FROM t2;
```

t1	t2
a	a
1	1
2	2
null	null

# NULL

```
SELECT * FROM t1  
UNION  
SELECT * FROM t2;
```

a
1
2
null

t1	t2
a	a
1	1
2	2
null	null

# NULL

More on NULLs:

- <http://www.xaprb.com/blog/2006/05/18/why-null-never-compares-false-to-anything-in-sql/>

# Advanced Subqueries

- **FROM** (*subquery* )
- **WHERE** *column\_name* < (*subquery*)
  - Can replace < with >, =, <=, >=, or <>
- **SELECT** (*subquery* )

# Aggregation

- Aggregate functions take a collection of values as input and return a single value.
- SQL offers five built-in aggregate functions:
  - Average: **AVG**
  - Minimum: **MIN**
  - Maximum: **MAX**
  - Total: **SUM**
  - Count: **COUNT**

# Aggregation

- Aggregate functions take a collection of values as input and return a single value.
- SQL offers five built-in aggregate functions:
  - Average: **AVG**
  - Minimum: **MIN**
  - Maximum: **MAX**
  - Total: **SUM**
  - Count: **COUNT**

The input to SUM and AVG must be collections of numbers, but the other operators can operate on collections of nonnumeric data types, such as strings, as well.

Except COUNT, all aggregate operations apply to a single attribute.

# Aggregation

- Retaining duplicates is important in computing **AVGs**, **SUMs**, and **COUNTs**.
- There are cases where we must eliminate duplicates before computing an aggregate function.
  - If we want to eliminate duplicates, we use the keyword **DISTINCT** in the aggregate expression.



# Aggregation

```
SELECT SUM(a)
FROM (VALUES (1), (1)) AS
T(a);
```

sum
2

```
SELECT SUM(DISTINCT a)
FROM (VALUES (1), (1)) AS
T(a);
```

sum
1

## Query 32

Calculate the average salary of all employees.

Note: Averages non-null salary values.

## Query 32

Calculate the average salary of all employees.

```
SELECT AVG(fsalary)  
FROM Faculty;
```

Note: Averages non-null salary values.

# Query 33

Display Faculty rows that earn more than the average salary.

## Query 33

Display Faculty rows that earn more than the average salary.

```
SELECT *  
FROM Faculty  
WHERE fsalary >  
      (SELECT AVG(fsalary)  
       FROM Faculty);
```

## Query 34

Find departments with faculty salaries over their budgeted amount.

## Query 34

Find departments with faculty salaries over their budgeted amount.

```
SELECT *  
FROM Department  
WHERE dsalary_budget <  
    (SELECT SUM(fsalary)  
     FROM Faculty  
     WHERE ddept = fdept);
```

# Aggregation with Grouping

- To apply the aggregate function not only to a single set of tuples, but also to a group of sets of tuples, use the **GROUP BY** clause.
  - The attribute or attributes given in the **GROUP BY** clause are used to form groups.
  - Tuples with the same value on all attributes in the **GROUP BY** clause are placed in one group.



# GROUP BY Syntax Issues

- All columns **SELECTed**
  - must be in **GROUP BY**  
or the target of an aggregate function.

# Query 35

Find the highest salary in each department.

## Query 35

Find the highest salary in each department.

```
SELECT fdept, MAX(fsalary)
FROM Faculty
GROUP BY fdept;
```

# GROUP BY Syntax Issues

- When a SQL query uses grouping, it is important to ensure that *the only attributes that appear in the **SELECT** statement without being aggregated are those that appear in the **GROUP BY** clause.*

# GROUP BY Syntax Issues

- Any attribute that is not present in the **GROUP BY** clause must appear only inside an aggregate function if it appears in the **SELECT** clause.

# The HAVING Clause

- At times, it is useful to state a condition that applies to groups rather than to tuples.
- SQL applies predicates in the **HAVING** clause after groups have been formed, so aggregate functions may be used.
- Any attribute that is present in the **HAVING** clause without being aggregated must appear in the **GROUP BY** clause.

# SQL Seduction 7

Find the highest salary in each department and include fid in the output.

# SQL Seduction 7

Find the highest salary in each department and include fid in the output.

```
SELECT fid, fdept, MAX(fsalary)
FROM Faculty
GROUP BY fdept;
```



# SQL Seduction 7

Find the highest salary in each department and include fid in the output.

```
SELECT fid, fdept, MAX(fsalary)  
FROM Faculty  
GROUP BY fdept;
```

**Syntax error! fid not in GROUP BY!**

# Query 36

This fixes the syntax error, but it doesn't solve the original problem.

```
SELECT fid, fdept, MAX(fsalary)
FROM Faculty
GROUP BY fdept, fid;
```



**WRONG result!**

**WHY?**

# Query 37

This solves the original problem.

```
SELECT fid, fdept, fsalary
FROM (SELECT fdept AS mdept, MAX(fsalary) AS msalary
      FROM Faculty
      GROUP BY fdept) AS t,
Faculty
WHERE (fdept = mdept)
AND fsalary = msalary;
```

Warning: The order of the “tables” matters in MySQL.

## Query 38

List all departments with only one course.

# Query 38

List all departments with only one course.

```
SELECT cdept FROM Course  
GROUP BY cdept  
HAVING count(*) = 1;
```

**HAVING** is to  
selecting groups  
as **WHERE** is to  
selecting rows.

# Aggregate Function Issues

- Can be **SELECT**ed
  - Typically used with **GROUP BY**
- Can appear in **HAVING** clause
- 
- Never allowed as a **WHERE** clause's simple condition, such as ~~**WHERE COUNT(\*) = 1**~~

# Sequence of Operations

- **WHERE** chooses rows
- **GROUP BY** groups chosen rows
- **HAVING** chooses groups
- **ORDER BY** sequences result
- **SELECT** chooses columns to display
- **DISTINCT** compresses duplicate result rows

# String Operations

- SQL specifies strings by enclosing them in single quotes.
  - A single quote character that is part of a string can be specified by using two single quote characters: 'Who's paying attention?'



# String Operations

- The SQL standard specifies that equality operations on strings is case sensitive, however some databases do not distinguish uppercase and lowercase while comparing strings.

**MySQL:**

**SELECT** 'foo' = 'FOO';

'foo' = 'FOO'
1

**PostgreSQL:**

**SELECT** 'foo' = 'FOO';

?column?
f

**SQLite:**

**SELECT** 'foo' = 'FOO';

0
---

# String Operations

<b>CONCAT</b>	Concatenation
<b>UPPER</b>	Converts string to uppercase.
<b>LOWER</b>	Converts string to lowercase.
<b>TRIM</b>	Removes spaces at the end of a string
<b>LENGTH</b>	Returns the length of the string.

\*\*\*This is just a small sample of the string functions. See the DB docs for many more.

- MySQL: <http://dev.mysql.com/doc/refman/5.7/en/string-functions.html>
- PostgreSQL: <https://www.postgresql.org/docs/9.5/static/functions-string.html>

# String Operations

**SELECT LENGTH('foo');**

?column?
3

**SELECT UPPER('foo');**

?column?
FOO

**SELECT LOWER('FOO');**

?column?
foo

**SELECT CONCAT('foo','bar');**

?column?
foobar

**SELECT CONCAT(TRIM('foo '), 'bar');**

?column?
foobar

# LIKE: Simple String Pattern Matching

- **s LIKE p**: pattern matching on strings
- **p** may contain two special symbols:
  - **%** = any sequence of characters
  - **\_** = any single character

# LIKE: Simple String Pattern Matching

SELECT 'foo' LIKE 'foo';

?column?
t

SELECT 'foo' LIKE '\_\_\_\_';

?column?
t

3 underscores



SELECT 'foo' LIKE 'fo';

?column?
f

SELECT 'foo' LIKE 'f\_';

?column?
f

# LIKE: Simple String Pattern Matching

```
SELECT 'foo' LIKE '%';
```

?column?
t

```
SELECT 'foo' LIKE CONCAT('%', 'f', '%');
```

?column?
t

```
SELECT 'foo' LIKE 'b%';
```

?column?
f

```
SELECT 'foo' LIKE CONCAT('%', x, '%')  
FROM (VALUES ('f'))  
AS t(x);
```

?column?
t

# Query 39

Find all the last names that begin with “B”.

## Query 39

Find all the last names that begin with “B”.

```
SELECT flast  
FROM Faculty  
WHERE flast LIKE 'B%';
```



# Query 40

Find all the last names that contain an “e”.

## Query 40

Find all the last names that contain an “e”.

```
SELECT flast  
FROM Faculty  
WHERE flast LIKE '%e%';
```

# Query 41

Find all the last names that start with “d” and have a “B” in the 4<sup>th</sup> position.

# Query 41

Find all the last names that start with “d” and have a “B” in the 4<sup>th</sup> position.

```
SELECT flast  
FROM Faculty  
WHERE flast LIKE 'd__B%';
```

## Query 42

List the names of employees who earn between \$20,000 and \$33,700 (inclusive).

# Query 42

List the names of employees who earn between \$20,000 and \$33,700 (inclusive)

```
SELECT flast, ffirst, fmi  
FROM Faculty  
WHERE fsalary BETWEEN 20000 AND 33700;
```

or

```
SELECT flast, ffirst, fmi  
FROM Faculty  
WHERE fsalary >= 20000 AND fsalary <= 33700;
```

# Temporal Math

- $\text{datetime} - \text{datetime} = \text{interval}$
- $\text{datetime} + \text{interval} = \text{datetime}$
- $\text{datetime} - \text{interval} = \text{datetime}$
- $\text{interval} + \text{interval} = \text{interval}$
- $\text{interval} - \text{interval} = \text{interval}$
- $\text{interval} + \text{numeric} = \text{interval}$
- $\text{interval} * \text{numeric} = \text{interval}$
- $\text{interval} / \text{numeric} = \text{interval}$

# Temporal Math

- *datetime* - *datetime* = *interval*

```
SELECT '20100110'::DATE - '20090110'::DATE;
```

?column?
365



# Temporal Math

- *datetime + interval = datetime*

```
SELECT '20100110'::DATE + 365;
```

?column?
2011-01-10

# Temporal Math

- *datetime - interval = datetime*

```
SELECT '20100110'::DATE - 365;
```

?column?
2009-01-10

# Pointless Complexity 1

```
SELECT DISTINCT fid  
FROM Department, Faculty;
```

reduces to

# Pointless Complexity 1

```
SELECT DISTINCT fid  
FROM Department, Faculty;
```

reduces to

```
SELECT fid FROM Faculty;
```

# Pointless Complexity 2

```
SELECT *  
FROM Faculty  
WHERE fid IN  
    (SELECT fid FROM Faculty);
```

reduces to

# Pointless Complexity 2

```
SELECT *  
FROM Faculty  
WHERE fid IN  
    (SELECT fid FROM Faculty);
```

reduces to

```
SELECT * FROM Faculty;
```

# Pointless Complexity 3

```
SELECT fid  
FROM Faculty  
GROUP BY fid  
HAVING COUNT(*) > 0;
```

reduces to

# Pointless Complexity 3

```
SELECT fid  
FROM Faculty  
GROUP BY fid  
HAVING COUNT(*) > 0;
```

reduces to

```
SELECT fid  
FROM Faculty;
```



# Pointless Complexity 4

```
SELECT ffirst  
FROM Faculty  
GROUP BY ffirst;
```

reduces to

# Pointless Complexity 4

```
SELECT ffirst  
FROM Faculty  
GROUP BY ffirst;
```

reduces to

```
SELECT DISTINCT ffirst  
FROM Faculty;
```

# SQL Seduction Summary

- **DISTINCT** compresses out duplicate rows
- Join tables based on their relationships
- Don't accidentally undo a correlation
- **NOT IN** can not be rewritten as  $\neq$
- One **IN** with 2 columns not the same as two **IN**s, one on each column
- Aggregate functions only look at non-null values
- Be mindful of **GROUP BY** syntax rules

As a reminder...

FACULTY						
fid	flast	ffirst	fmi	fdept	fsalary	fmgr_id
12058	Borys	Ted	J	CSI	48000	22321
12206	Ryan	Alfred	C	ENG	48000	52110
21004	Perry	Bill	S	BIO	21800	31890
22321	Brady	Kathy	M	CSI	63400	52110
31890	Coulsen	Mary	null	BIO	21400	52110
32000	delBene	Bill	S	CSI	63500	22321
47862	Anders	John	P	ENG	33700	12206
52110	Smith	Alice	null	ADM	82000	null

# Query 43

Display all the faculty columns with **fmi** equal to “S”.

```
SELECT *  
FROM Faculty  
WHERE fmi = “S”;
```

fid	flast	ffirst	fmi	fdept	fsalary	fmgr_id
21004	Perry	Bill	S	BIO	21800	31890
32000	delBene	Bill	S	CSI	63500	22321

# Query 44

Display all the faculty columns with **fmi** not equal to “S”

```
SELECT *  
FROM Faculty  
WHERE fmi <> “S”;
```

fid	flast	ffirst	fmi	fdept	fsalary	fmgr_id
12058	Borys	Ted	J	CSI	48000	22321
12206	Ryan	Alfred	C	ENG	48000	52110
22321	Brady	Kathy	M	CSI	63400	52110
47862	Anders	John	P	ENG	33700	12206

# Query 45

Combine all the rows from Query 43 and 44

```
SELECT * FROM Faculty WHERE fmi = "S"  
UNION ALL  
SELECT * FROM Faculty WHERE fmi <> "S"
```

Only 6 rows returned, not 8.

## Query 46

Display all the faculty columns with unknown **fmi**.

```
SELECT *  
FROM Faculty  
WHERE fmi IS NULL;
```



# Query 47

Combine all the rows from Query 43, 44, and 46.

```
SELECT * FROM Faculty WHERE fmi = "S"  
UNION ALL  
SELECT * FROM Faculty WHERE fmi <> "S"  
UNION ALL  
SELECT * FROM Faculty WHERE fmi IS NULL;
```

## Query 48

Run this query to see why queries 43 through 47 work the way they do.

```
SELECT fmi,  
       fmi = 'S' AS "=S",  
       fmi <> 'S' AS "<> S",  
       (fmi IS NULL) AS "IS NULL",  
       CHAR_LENGTH(fmi) AS length  
FROM Faculty;
```

Note that true is 1 and false is 0.

Syntax valid for both Postgres and MySQL.

# Query 49

Find faculty where fmi appears in fmi column in faculty.

```
SELECT *  
FROM Faculty  
WHERE fmi IN  
    (SELECT fmi FROM Faculty);
```

Only 6 rows appear in the result.

## Query 50

Find faculty where fmi doesn't appear in fmi column in faculty.

```
SELECT *  
FROM Faculty  
WHERE fmi NOT IN  
  (SELECT fmi FROM Faculty);
```

No rows appear in the result.

# Query 51

Find faculty where fmi appears in fmi column in faculty using **EXISTS** construct.

```
SELECT * FROM Faculty AS a
WHERE EXISTS
  (SELECT b.fmi
   FROM Faculty AS b
   WHERE a.fid = b.fid);
```

All 8 rows appear in the result – never get empty set.

# Query 51

Find faculty where fmi appears in fmi column in faculty using **EXISTS** construct.

```
SELECT * FROM Faculty AS a
WHERE EXISTS
  (SELECT b.fmi
   FROM Faculty AS b
   WHERE a.fid = b.fid);
```

Primary keys are guaranteed to never be null. Therefore, this condition must evaluate to true for one row. The result of the inner query may be a row with a null for its only attribute, but it is a row nonetheless.

All 8 rows appear in the result – never get empty set.

# Query 52

Find faculty where fmi doesn't appear in fmi column in faculty using **EXISTS** construct.

```
SELECT *  
FROM Faculty AS a  
WHERE NOT EXISTS  
  (SELECT b.fmi  
   FROM Faculty AS b  
   WHERE a.fid = b.fid);
```

No rows appear  
in the result.

# Query 53

Count the number of courses offered by each department, and include zero counts.



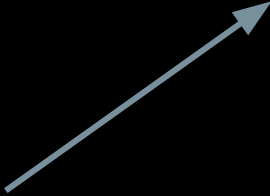
# Query 53

Count the number of courses offered by each department, and include zero counts.

```
SELECT ddept, dname, COUNT(*) AS count
FROM Department, Course
WHERE ddept = cdept
GROUP BY ddept, dname;
```

ddept	dname	count
ATM	Atmospheric Science	2
BIO	Biology	2
CSI	Computer Science	3
ENG	English	1

Does not include  
the zero counts!



# Query 53

Count the number of courses offered by each department, and include zero counts.

```
SELECT ddept, dname, COUNT(*) AS count
FROM Department, Course
WHERE ddept = cdept
GROUP BY ddept, dname
UNION
SELECT ddept, dname, 0
FROM Department WHERE ddept NOT IN
(SELECT cdept FROM Course);
```

ddept	dname	count
ADM	Administration	0
ATM	Atmospheric Science	2
BIO	Biology	2
CSI	Computer Science	3
ENG	English	1
SPN	Spanish	0

# Query 54

Rewrite Query 53 without the **UNION**.

# Query 54

Rewrite Query 53 without the **UNION**.

```
SELECT ddept,  
       dname,  
       (SELECT COUNT(*)  
        FROM Course  
        WHERE ddept = cdept) AS count  
FROM Department;
```

## Query 55

Without using **LIMIT**, find the 3 highest salaries in faculty.

## Query 55

Without using **LIMIT**, find the 3 highest salaries in faculty.

```
SELECT COUNT(*), a.fsalary  
FROM Faculty AS a, Faculty AS b  
WHERE a.fsalary <= b.fsalary  
GROUP BY a.fsalary  
HAVING COUNT(*) <= 3;
```

# Query 55a

Consider:

```
SELECT a.fid, a.fsalary, b.fid, b.fsalary  
FROM Faculty AS a, Faculty AS b  
WHERE a.fsalary <= b.fsalary  
ORDER BY a.fsalary DESC, b.fsalary DESC;
```

# Query 55a Results

a.fid	a.fsalary	b.fid	b.fsalary
52110	\$82,000.00	52110	\$82,000.00
32000	\$63,500.00	52110	\$82,000.00
32000	\$63,500.00	32000	\$63,500.00
22321	\$63,400.00	52110	\$82,000.00
22321	\$63,400.00	32000	\$63,500.00
22321	\$63,400.00	22321	\$63,400.00
12206	\$48,000.00	52110	\$82,000.00
12058	\$48,000.00	52110	\$82,000.00
12206	\$48,000.00	32000	\$63,500.00
12058	\$48,000.00	32000	\$63,500.00
12058	\$48,000.00	22321	\$63,400.00
12206	\$48,000.00	22321	\$63,400.00
12058	\$48,000.00	12206	\$48,000.00
12206	\$48,000.00	12058	\$48,000.00
12058	\$48,000.00	12058	\$48,000.00
12206	\$48,000.00	12206	\$48,000.00
47862	\$33,700.00	52110	\$82,000.00

1  
2  
3

10



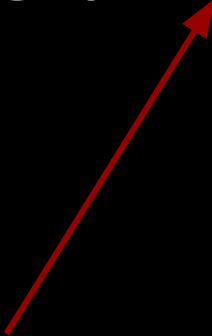
## Query 56

Find the 6 highest salaries in faculty using the same strategy as Query 55.

# Query 56

Find the 6 highest salaries in faculty using the same strategy as Query 55.

```
SELECT count(*), a.fsalary
FROM Faculty AS a, Faculty AS b
WHERE a.fsalary <= b.fsalary
GROUP BY a.fsalary
HAVING COUNT(*) <= 6;
```



count(*)	fsalary
6	33700
3	63400
2	63500
1	82000

Duplicate salary values highlight a problem.

# Query 56

Duplicate salary values highlight a problem.

```
SELECT fsalary, COUNT(*) as ct
FROM Faculty
GROUP BY fsalary
ORDER BY ct DESC, fsalary DESC;
```

fsalary	count(*)
48000	2
82000	1
63500	1
63400	1
33700	1
21800	1
21400	1

# Query 56 Results

Query56 - Microsoft Access

Home Create External Data Database Tools Acrobat

View Clipboard Font Rich Text Records Filter Sort & Filter Size to Fit Form Switch Windows Find

All Access Objects

Query43  
Query44  
Query46  
Query48  
Query49  
Query50  
Query51  
Query52  
Query54  
Query55  
Query55a  
Query56  
Query57  
Query58  
Query59

count	fsalary
6	\$33,700.00
3	\$63,400.00
2	\$63,500.00
1	\$82,000.00

Record: 1 of 4 No Filter Search

Datasheet View Num Lock SQL

# More of Query 55a Results

Query55a - Microsoft Access

Home Create External Data Database Tools Acrobat

View Clipboard Font Rich Text Records Filter Size to Fit Form Switch Windows Find

All Access Objects

Query43 Query44 Query46 Query48 Query49 Query50 Query51 Query52 Query54 Query55 Query55a Query56 Query57 Query58 Query59

a.fid	a.fsalary	b.fid	b.fsalary
47862	\$33,700.00	52110	\$82,000.00
47862	\$33,700.00	32000	\$63,500.00
47862	\$33,700.00	22321	\$63,400.00
47862	\$33,700.00	12206	\$48,000.00
47862	\$33,700.00	12058	\$48,000.00
47862	\$33,700.00	47862	\$33,700.00
21004	\$21,800.00	52110	\$82,000.00
21004	\$21,800.00	32000	\$63,500.00
21004	\$21,800.00	22321	\$63,400.00
21004	\$21,800.00	12058	\$48,000.00
21004	\$21,800.00	12206	\$48,000.00
21004	\$21,800.00	47862	\$33,700.00
21004	\$21,800.00	21004	\$21,800.00
31890	\$21,400.00	52110	\$82,000.00
31890	\$21,400.00	32000	\$63,500.00
31890	\$21,400.00	22321	\$63,400.00
31890	\$21,400.00	12058	\$48,000.00

Record: 17 of 37 No Filter Search

Datasheet View Num Lock SQL

# Query 57

Better way to find the 6 highest salaries in faculty.

```
SELECT COUNT(*), a.fsalary
FROM
  (SELECT DISTINCT a.fsalary, b.fsalary
   FROM Faculty AS a, Faculty AS b
   WHERE a.fsalary <= b.fsalary)
GROUP BY a.fsalary
HAVING COUNT(*) <= 6;
```

# Query 57 Results

Query57 - Microsoft Access

Home Create External Data Database Tools Acrobat

View Clipboard Font Rich Text Records Filter Sort & Filter Size to Fit Form Switch Windows Find

All Access Objects

Query43  
Query44  
Query46  
Query48  
Query49  
Query50  
Query51  
Query52  
Query54  
Query55  
Query55a  
Query56  
Query57  
Query58  
Query59

count	fsalary
6	\$21,800.00
5	\$33,700.00
4	\$48,000.00
3	\$63,400.00
2	\$63,500.00
1	\$82,000.00

Record: 1 of 6 No Filter Search

Datasheet View Num Lock SQL

# Query 57a

Consider:

```
SELECT DISTINCT a.fsalary, b.fsalary  
FROM Faculty AS a, Faculty AS b  
WHERE a.fsalary <= b.fsalary  
ORDER BY a.fsalary DESC, b.fsalary DESC;
```



# Query 57a Results

Query57a - Microsoft Access

Home Create External Data Database Tools Acrobat

View Clipboard Font Rich Text Records Filter Sort & Filter Size to Fit Form Switch Windows Find

All Access Objects

	a.fsalary	b.fsalary
Query50	\$82,000.00	\$82,000.00
Query51	\$63,500.00	\$82,000.00
Query52	\$63,500.00	\$63,500.00
Query54	\$63,400.00	\$82,000.00
Query55	\$63,400.00	\$63,500.00
Query55a	\$63,400.00	\$63,400.00
Query56	\$48,000.00	\$82,000.00
Query57	\$48,000.00	\$63,500.00
Query57a	\$48,000.00	\$63,400.00
Query58	\$48,000.00	\$48,000.00
Query59	\$33,700.00	\$82,000.00
Query60	\$33,700.00	\$63,500.00
Query61	\$33,700.00	\$63,400.00
SQL Seduction 1	\$33,700.00	\$48,000.00
SQL Seduction 3	\$33,700.00	\$33,700.00
	\$21,800.00	\$82,000.00
	\$21,800.00	\$63,500.00

Record: 1 of 28 No Filter Search

Ready Num Lock SQL

# Join Types

- **INNER JOIN**
  - Data must match in both tables
- **LEFT JOIN** or **LEFT OUTER JOIN**
  - Data must match in both tables or appear in left table
- **RIGHT JOIN** or **RIGHT OUTER JOIN**
  - Data must match in both tables or appear in right table

# Query 58

Original way to do an inner join

```
SELECT *  
FROM Department, Faculty  
WHERE ddept = fdept;
```

# Query 59

Syntax introduced in SQL92 standard.

```
SELECT *  
FROM Department INNER JOIN Faculty  
ON ddept = fdept;
```

# Query 60

Three-way table join:

```
SELECT *  
FROM Department  
  INNER JOIN Faculty  
    ON ddept = fdept  
  INNER JOIN Section  
    ON fid = sid;
```

# Query 61

Join department and faculty, and include departments with no faculty.

# Query 61

Join department and faculty, and include departments with no faculty.

```
SELECT *  
FROM Department  
    LEFT JOIN Faculty  
        ON ddept = fdept;
```

# Query 62

Rewrite Query 61 as a RIGHT JOIN

```
SELECT *  
FROM Faculty  
    RIGHT JOIN Department  
    ON ddept = fdept;
```



# Full Outer Join

- Combine LEFT and RIGHT joins.

# Full Outer Join

- MySQL

SELECT \*

FROM t1 LEFT JOIN t2 ON t1.a = t2.a

UNION ALL

SELECT \*

FROM t1 RIGHT JOIN t2 ON t1.a = t2.a;

a	a
1	1
2	null
null	null
1	1
null	null
null	3

t1	t2
a	a
1	1
2	null
null	3

# Full Outer Join

- MySQL

SELECT \*

FROM t1 LEFT JOIN t2 ON t1.a = t2.a

UNION

SELECT \*

FROM t1 RIGHT JOIN t2 ON t1.a = t2.a;

a	a
1	1
2	null
null	null
null	3

t1	t2
a	a
1	1
2	null
null	3

# Full Outer Join

- PostgreSQL

SELECT \*

FROM t1 FULL OUTER JOIN t2

ON t1.a = t2.a ;

a	a
1	1
2	null
null	null
null	3
null	null

t1	t2
a	a
1	1
2	null
null	3

# Cautions

- Watch out for `NULL` values
- Subqueries can be used in many places
  - `IN`
  - `EXISTS`
  - `FROM`
  - `WHERE`
  - `SELECT`
- `JOIN` syntax introduced in SQL92 standard
- `IN` subquery syntax added in SQL99